

Two dimensionless ratios define operational limits in hierarchical-memory inference

Geunsik Lim ^{1*}

¹ College of Information and Communication Engineering,
Sungkyunkwan University (SKKU), Suwon, South Korea .

Corresponding author(s). E-mail(s): leemgs@g.skku.edu;

Abstract

We formulate hierarchical-memory inference using two dimensionless ratios: fast-memory residency $R_C = C_{\text{fast}}/W$ and compute-transfer overlap $R_B = T_{\text{comp}}/T_{\text{transfer}}$. These ratios define a capacity-limited region ($R_C < 0.5$), an I/O-limited region ($R_C \geq 0.5$, $R_B < 1$), and a coordination-dominated region ($R_C \geq 0.5$, $R_B \geq 1$). We derive irreducible capacity and transfer-cost bounds and release a measurement harness using wall-clock or device-event timing. A compiled dependent-load microbenchmark on one Intel Xeon CPU shows mean access time increasing from 6.8 ns for a 0.5 MiB working set to 169.7 ns for 256 MiB. The increase spans several cache and translation levels and does not identify a discontinuity at $R_C = 0.5$. On 3 accelerators of two vendors (NVIDIA Tesla T4, NVIDIA A100, and Google TPU v5e), the residency prediction holds—per-step latency rises as residency falls below $R_C = 0.5$ on every device (1.23–2.83 $\times$)—and on the Tesla T4 the I/O-limited label appears exactly where the derived overlap boundary $R_B < 1$ predicts. These results establish the two-ratio labels as a measurement discipline across real silicon; per-device boundary values and production-workload generality remain open.

Keywords: hierarchical memory systems, memory orchestration, AI inference, working-set residency, compute-transfer overlap, reproducible measurement

1 Introduction

Modern foundation models have shifted the performance frontier of AI inference from arithmetic throughput to *data coordination*. When model weights, activations, and key-value (KV) caches exceed accelerator memory, inference latency is governed

less by compute and more by how efficiently execution coordinates data movement across a hierarchical memory system spanning device memory, host memory, and secondary storage [1, 2]. Because inference now dominates the recurring cost and energy of deployed models, this coordination has direct economic and environmental stakes [3–5]. Memory-bound inference spans autoregressive language models [6, 7], vision and vision–language models [8, 9], and retrieval-augmented pipelines [10], so the coordination question recurs across otherwise different workloads.

Hardware analyses identify memory capacity, bandwidth, and interconnect latency as important constraints on inference [11–13]. Separating capacity pressure from exposed transfer is therefore useful when reporting and comparing inference systems.

Existing systems reduce memory pressure along several axes. Serving frameworks manage the KV cache and schedule requests to raise throughput [14–17]; offloading and memory-pooling systems spill weights or activations to host DRAM or NVMe [18–23]; swapping systems stage tensors on and off the accelerator [24–26]; and compression—sparsity, pruning, and IO-aware attention—shrinks the working set or its traffic [27–29]. Parallelism redistributes memory across devices [30, 31], while host-side allocators and NUMA-aware placement shape the sustained transfer path [32–35]. The design space is large enough to warrant several recent surveys [36, 37], and it widens further in edge and heterogeneous deployments [38–41]. Across these approaches the same technique can help or hurt depending on model size, batch size, hardware topology, and how much compulsory traffic remains on the critical path. This motivates a compact description that exposes those conditions before comparing policies.

We take a regime-based view. Denning’s working-set model and thrashing theory [42] identify capacity saturation as a binary event on a *single-tier* system, and the Roofline model [43] expresses performance as $\min(F \cdot I, B_{\text{peak}})$ for a fixed workload on a single tier. Neither framework addresses a *multi-tier* hierarchy in which computation, locality, cross-tier transfer, and synchronisation co-exist as distinct, adjustable latency terms. Building instead on classical queueing analysis of overlapping services [44], our accounting model decomposes end-to-end latency into four terms ($T_{\text{comp}} + T_{\text{mem}} + T_{\text{swap}} + T_{\text{sync}}$) and uses two dimensionless ratios to label an operating point. The three operational labels are governed by the fast-memory residency ratio $R_C = C_{\text{fast}}/W$ and the overlap ratio $R_B = B_{\text{slow}}T_{\text{comp}}/D = T_{\text{comp}}/T_{\text{transfer}}$. The labels state whether majority residency is available and whether ideal compute–transfer overlap can hide compulsory traffic; they do not by themselves predict a policy’s speed-up.

We make three contributions:

- **Operational formulation.** We separate residency from compute–transfer overlap using two measurable, dimensionless ratios and state an explicit three-label classification rule.
- **Minimal analytical model.** A structural lower bound on the irreducible memory-access and transfer costs, with regime boundaries following directly from the two ratios: $\theta_B = 1.0$ derived as the overlap boundary ($R_B = T_{\text{comp}}/T_{\text{transfer}} = 1$), and $\theta_C = 0.50$ as the majority-residency crossover.

• **Open proof of concept.** A compiled dependent-load CPU probe illustrates working-set sensitivity, and a layered-matrix harness runs on a CPU and on 3 accelerators of two vendors (NVIDIA Tesla T4, NVIDIA A100, Google TPU v5e). The residency prediction holds on all three accelerators and the overlap boundary is resolved on the Tesla T4. We distinguish this proof of concept from a preregistered multi-platform campaign.

Together, the formulation and measurements provide a reproducible starting point for testing when hierarchical-memory coordination is constrained by residency or exposed transfer. They confirm the residency direction on one CPU and three accelerators, and the overlap boundary on one GPU, but do not establish effect sizes, strategy rankings, or per-device boundary values.

2 Results

2.1 A two-ratio operational model

We describe a hierarchical-memory operating point using

$$R_C = \frac{C_{\text{fast}}}{W}, \quad R_B = \frac{B_{\text{slow}} T_{\text{comp}}}{D} = \frac{T_{\text{comp}}}{T_{\text{transfer}}}, \quad (1)$$

where C_{fast} is usable fast-memory capacity, W is the active working set, D is compulsory cross-tier traffic per step, B_{slow} is sustained bandwidth on the limiting link, and $T_{\text{transfer}} = D/B_{\text{slow}}$. Unlike a ratio that uses total step time, R_B is not circular: it compares independently timed computation and transfer intervals.

We use $\theta_C = 0.50$ as an explicit majority-residency convention and derive $\theta_B = 1$ from equality of computation and transfer time. These definitions produce three operational labels:

- **capacity-limited:** $R_C < \theta_C$;
- **I/O-limited:** $R_C \geq \theta_C$ and $R_B < \theta_B$; and
- **coordination-dominated:** $R_C \geq \theta_C$ and $R_B \geq \theta_B$.

The capacity boundary is a convention, not a thermodynamic critical point. The I/O boundary is an overlap boundary: for $R_B < 1$, at least $T_{\text{transfer}} - T_{\text{comp}}$ cannot be hidden by ideal overlap.

2.2 Irreducible costs and falsifiable predictions

For accounting purposes, define non-overlapping critical-path contributions so that a step can be written as

$$T_{\text{total}} = T_{\text{comp}} + T_{\text{mem}} + T_{\text{swap}} + T_{\text{sync}}. \quad (2)$$




Fig. 1 Analytical regime map. The vertical line is the declared majority-residency convention $\theta_C = 0.50$; the horizontal segment is the derived overlap boundary $\theta_B = 1.0$. The diagram is generated from the same constants as the released classifier and contains no empirical data.

103 Independently of that accounting decomposition, a conservative service-time lower
 104 bound that allows the underlying services to overlap is

$$T_{\text{total}} \geq \max \left\{ T_{\text{comp}}, \rho W \max(0, 1 - R_C), \frac{D}{B_{\text{slow}}} \right\}, \quad (3)$$

105 where W denotes bytes accessed in the step and ρ is a lower bound on service time
 106 per non-resident byte for the specified access pattern. Taking the maximum permits
 107 ideal overlap and avoids double-counting bytes that contribute to both a cache miss
 108 and a cross-tier transfer. Synchronisation and contention can only increase observed
 109 time, so Equation (3) is not a performance predictor. The residual above it cannot be
 110 assigned to one latency component without event-level measurements.

111 The formulation makes two falsifiable predictions. First, reducing residency while
 112 holding the access pattern fixed should expose a higher memory-access cost. Second,
 113 when $R_C \geq \theta_C$, configurations with $R_B < 1$ must expose transfer time even under
 114 ideal compute-transfer overlap. It does not, without additional measurements, predict
 115 the magnitude of a strategy's speed-up or claim that strategy rankings necessarily
 116 invert.

Table 1 Selected measured points from the compiled CPU pointer-chain experiment. Values are mean $\pm$ one standard deviation over seven trials.

Working set (MiB)	R_C	Latency (ns/load)	Interpretation
0.5	96.000	6.8 ± 0.1	small working set
48	1.000	118.5 ± 6.6	reported LLC capacity
96	0.500	132.5 ± 14.3	residency convention
256	0.188	169.7 ± 10.4	largest working set

2.3 CPU proof-of-concept measurements

We evaluated working-set sensitivity on an Intel Xeon Platinum 8370C CPU exposed through a virtualised Linux environment. The experiment used a randomised pointer chain over working sets from 0.5 to 256 MiB and seven trials per point. A compiled C loop performed dependent loads and timed them with `CLOCK_MONOTONIC_RAW`; Python was used only to construct the chain and summarise trials. Linux sysfs reported a 48 MiB last-level cache, which we used as C_{fast} . Mean dependent-load time was 6.8 ± 0.1 ns at 0.5 MiB ($R_C = 96$), 118.5 ± 6.6 ns at 48 MiB ($R_C = 1$), 132.5 ± 14.3 ns at 96 MiB ($R_C = 0.5$), and 169.7 ± 10.4 ns at 256 MiB ($R_C = 0.188$), where uncertainty is one standard deviation across trials. The non-monotonic intermediate points show that cache level, translation, virtualisation, and system noise matter. In particular, these data do not show a discontinuity at the conventional $R_C = 0.5$ label.

As a separate harness check, we used the layered matrix program with 12 layers, $d = 512$, batch size 8, and five timing windows. It varied the fraction of layer weights retained in the fast tier. The mean total latency averaged over the three sampled points below $R_C = 0.5$ was 4.03 ms, compared with 1.87 ms for the fully resident point, a 2.15-fold descriptive ratio. Because computation and copying share the same CPU memory system, the accompanying R_B sweep was noisy and is not used as evidence for the I/O boundary.

These measurements are proof-of-concept results from one CPU, not a substitute for the previously planned A100, MI250, TPU v4, Inferentia2, and Optane campaign. Accordingly, we make no quantitative cross-accelerator, classifier- accuracy, energy, or strategy-inversion claim in this version.

2.4 Accelerator proof-of-concept measurements

We ran the layered program on 3 real accelerators of distinct architectures and two vendors—an NVIDIA Tesla T4, an NVIDIA A100, and a Google TPU v5e—with $d = 2048$, 24 layers, batch size 8, and ten timing windows. On the CUDA devices, computation and host-to-device transfer were timed separately with CUDA events; non-resident layer weights were staged in pinned host memory and copied per step, so residency (R_C) and overlap (R_B) were varied by construction. These runs test the two falsifiable predictions of Section 2.2 on hardware whose device–host link, unlike a single CPU’s memory system, makes the exposed-transfer regime physically reachable.

Prediction 1 (residency) holds on all 3 accelerators. On every device, mean per-step latency in the capacity-limited region ($R_C < \theta_C$) exceeded that at full residency

Table 2 Capacity crossover across 3 accelerators (layered probe, $d = 2048$). T is mean per-step latency; “Ratio” is capacity-limited over resident. The direction (ratio > 1) is the predicted effect; magnitudes are device-dependent and not predicted.

Accelerator	Resident T (ms)	Cap.-limited T (ms)	Ratio	Overlap sweep
Tesla T4	14.5	41.1	$2.83\times$	overlap boundary resolved
NVIDIA A100	52.4	70.6	$1.35\times$	sweep B compute-bound
Google TPU v5e	220.6	270.6	$1.23\times$	R_B split unreliable (XLA)

($R_C \geq 1$), by 1.23 – $2.83\times$ (Table 2). The formulation predicts this *direction*, not its magnitude; the magnitude shrinks on higher-bandwidth devices, where the offload penalty is smaller relative to compute. Near $R_C = 1$ the faster devices show non-monotonic noise, consistent with the small per-step workload ($d = 2048$) leaving launch and measurement overhead visible.

Prediction 2 (overlap) is resolved on the T4. Holding residency fixed and scaling compute moved the T4 operating point across the derived boundary $\theta_B = 1$: every point with $R_B < 1$ was labelled I/O-limited and every point with $R_B \geq 1$ coordination-dominated, with the transition between $R_B = 0.90$ and $R_B = 1.19$ —not at a fitted value. On the A100 the controlled sweep stayed compute-bound ($R_B > 2$) and so did not probe the boundary, and on the TPU the XLA path’s deferred execution makes the per-op R_B split unreliable; both need a larger-workload, preregistered run. The overlap boundary is therefore established on one GPU, while the capacity direction is established across three accelerators.

These are proof-of-concept results; they confirm the residency prediction across architectures and the overlap boundary on one GPU, but do not establish per-device boundary values, strategy inversion, classifier accuracy, or energy claims, which still require a larger, preregistered multi-platform campaign.

2.5 Reproducibility boundary

The repository contains the raw JSONL records, machine-readable summaries, and scripts used for Table 1. It also contains a simulated backend for testing the analysis pipeline. Simulated output is labelled as such and is not used as empirical evidence. Accelerator support in the measurement harness is an executable protocol for future validation; it is not evidence that those accelerators were measured in this study.

3 Discussion

The two-ratio description separates two questions that are often conflated in hierarchical-memory optimisation: whether the active data can reside in the fast tier, and whether compulsory transfer can be hidden behind computation. The resulting labels are operational rather than universal phases. $\theta_B = 1$ follows exactly from the

overlap definition, whereas $\theta_C = 0.5$ is a declared midpoint that makes “majority resident” precise. A different application may choose another residency threshold without changing Equation (3).

The CPU probes show that latency depends strongly on working-set residency, but they do not validate a sharp boundary at $R_C = 0.5$. They stress different mechanisms: dependent pointer loads probe cache and translation effects, while the layered matrix harness combines weight copying with computation. The accelerator runs (Section 2.4) exercise real DMA and asynchronous CUDA streams across three devices of two vendors. The residency prediction holds on all three; the overlap crossover at $\theta_B = 1$ is resolved on the T4, while the A100 sweep stayed compute-bound and the TPU’s XLA timing is unreliable. The data do not establish per-device boundary *values*, HBM- or collective-communication effects at scale, or the generality of the three labels across model families.

The framework is most useful as a measurement discipline. Reporting R_C and R_B alongside latency makes hidden assumptions about residency and overlap inspectable. It also prevents a common circularity: using total step duration in both the numerator of a bandwidth ratio and as the outcome being explained. For deployment, C_{fast} , W , D , sustained bandwidth, and isolated compute time must be measured for the actual workload and hardware rather than copied from nominal data-sheet values.

Several limitations are material. The present evidence comes from one CPU and three accelerators (two GPUs and a TPU) with small synthetic probes; it does not include production models, concurrent requests, accelerator energy measurements, strategy comparisons, or unseen-hardware classifier tests. The standard deviations describe repeated trials on one machine and are not confidence intervals over a population of platforms. These measurements are exploratory technical checks rather than independent hardware replicates, and no causal inference is supported. The pointer-chain endpoints also span several hierarchy transitions, so their ratio should not be interpreted as a universal effect size. Finally, Equation (2) is an accounting identity only when its terms are defined as non-overlapping critical-path contributions; the max-form lower bound avoids assuming that decomposition in measured systems.

A decisive follow-up should preregister operating-point grids around $R_C = 0.5$ and $R_B = 1$, collect raw event traces on at least two accelerator architectures, and compare fixed orchestration policies without changing other workload parameters. Such data could test boundary sharpness, policy ranking, and generalisation. Until those measurements are available, the contribution of this work is the corrected analytical formulation, the open protocol, and a CPU-and-multi-accelerator proof of concept rather than a claim of universal accelerator validation.

4 Methods

Definitions and classification

We calculate $R_C = C_{\text{fast}}/W$ from a declared usable fast-tier capacity and active working-set size. We calculate $R_B = T_{\text{comp}}/T_{\text{transfer}}$, with compute and transfer timed separately. Classification gives capacity precedence: $R_C < 0.5$ is capacity-limited;

222 otherwise $R_B < 1$ is I/O-limited; all remaining points are coordination-dominated.
223 This rule is implemented in the released Python code.

224 **Pointer-chain capacity probe**

225 The CPU probe allocates a NumPy array and constructs a seeded random permutation.
226 A C11 kernel compiled with `-O3` traverses it as a dependent pointer chain;
227 each loaded index determines the address of the next load, limiting vectorisation,
228 instruction-level parallelism, and prefetching. We measured working sets of 0.5, 1, 2, 4,
229 8, 16, 24, 32, 48, 64, 96, 128, 192, and 256 MiB. Each point used seven timed trials after
230 warm-up. The output stores working-set size, R_C , mean nanoseconds per dependent
231 load, standard deviation, and trial count. The fast-tier capacity was read from Linux
232 `sysfs` and was 48 MiB. The host exposed three logical CPUs of an Intel Xeon Platinum
233 8370C under KVM; the summary records the kernel, Python, and NumPy versions.
234 Timing used `clock_gettime(CLOCK_MONOTONIC_RAW)` inside the compiled loop. No
235 outlier was removed and CPU affinity or frequency was not controlled.

236 **Layered matrix and copy probe**

237 The second probe allocates square, float32 layer matrices; a selected fraction remains
238 resident and non-resident matrices are copied before multiplication. Matrix entries
239 are deterministically seeded and scaled by $1/\sqrt{d}$ to prevent numerical overflow during
240 repeated multiplication. On the CPU we used NumPy’s wall-clock path with 12 layers,
241 $d = 512$, and five windows per point. On the accelerators (Section 2.4) the same
242 program ran with 24 layers, $d = 2048$, batch size 8, ten windows, and pinned host
243 staging: on the NVIDIA Tesla T4 and A100 through the harness’s CUDA-event path,
244 timing computation and host-to-device transfer separately with `torch.cuda.Event`;
245 and on a Google TPU v5e through the `torch.xla` path. Because XLA executes lazily,
246 per-op R_B timing on the TPU is unreliable and we treat its R_B split as indicative
247 only, relying on end-to-end latency for the residency test.

248 For the reported capacity ratio, we averaged total latency at the three sampled
249 points below $R_C = 0.5$ and divided by latency at $R_C = 1$. This endpoint summary is
250 descriptive; no hypothesis test or population-level confidence interval is claimed. The
251 pointer-chain table reports mean and population standard deviation across the seven
252 repeated trials exactly as stored in the JSON summary.

253 **Software and provenance**

254 The compiled pointer-chain run used Python 3.12.13, NumPy 2.5.2, and GCC 13.3;
255 the earlier layered-matrix harness check used Python 3.11.15, but that summary did
256 not capture the NumPy version. Seeds, grids, timing code, raw records, summaries,
257 and commands are version controlled. The analytical regime map reads threshold
258 constants from the same module as the classifier. The simulated backend is provided
259 solely for software testing and qualitative exploration; its outputs are excluded from
260 the reported measurements.

261 Statistics and reporting

262 This proof-of-concept study did not use random assignment, blinding, or a sample-size
263 calculation. Trial counts were fixed before summarisation. We report all measured grid
264 points in the machine-readable data, selected points in Table 1, arithmetic means, and
265 one standard deviation. No claim of statistical significance is made. Hardware-platform
266 generalisation requires independent machines and is left for future work.

267 Data availability

268 All data supporting the reported CPU measurements are included in the pub-
269 lic repository at <https://github.com/leemgs/orion>, under `code/results/cpu_probe/`
270 and `code/results/colab_probe/`. These directories contain the raw JSONL records
271 and machine-readable JSON summaries used in the text and table. No accelerator
272 data are reported in this article.

273 Code availability

274 The analysis and measurement code is available at <https://github.com/leemgs/orion>.
275 The `code/experiments/` directory contains the CPU probes, analytical-figure gener-
276 ator, and optional CUDA/XLA measurement paths. Exact reproduction commands
277 are provided in the README and Supplementary Information. A separately labelled
278 simulated backend is supplied for software testing; simulated output is not used as
279 evidence in this article.

280 Author Contributions

281 G.L.: Conceptualization, Formal analysis, Investigation, Methodology, Software, Val-
282 idation, Visualization, Writing – original draft, Writing – review & editing.

283 Competing Interests

284 The author declares no competing interests.

285 Acknowledgements

286 We thank Donghyeon Kang, Seungmin Oh and Yunsu Lee for valuable feedback. We
287 are also grateful to Wooyeon Lee, Hyoyoung Kim, and Heekuk Lee, experts in systems
288 and memory, whose insights and meaningful comments greatly improved this work.

289 References

- 290 [1] Xiaoyu Ma and David Patterson. Challenges and research directions for large
291 language model inference hardware. *IEEE Computer*, 59, 2026. URL <https://arxiv.org/abs/2601.05047>. Accepted.
292

- [2] Amir Gholami, Zhewei Yao, Sehoon Kim, Coleman Hooper, Michael W. Mahoney, and Kurt Keutzer. Ai and memory wall. *IEEE Micro*, 2024. URL <https://arxiv.org/abs/2403.14123>.
- [3] International Energy Agency. Electricity 2024: Analysis and forecast to 2026. Technical report, IEA, 2024. URL <https://www.iea.org/reports/electricity-2024>. Global average carbon intensity: 494 gCO₂/kWh.
- [4] Jovan Stojkovic, Chaojie Zhang, Íñigo Goiri, Josep Torrellas, and Esha Choukse. DynamoLLM: Designing LLM inference clusters for performance and energy efficiency. In *Proc. HPCA*, 2025. URL <https://arxiv.org/abs/2408.00741>.
- [5] Isabel Juniewicz. Hyperscaler capex is on trend to outpace their cash inflows by the end of 2026. Epoch AI, Data Insight, 2026. Accessed: 2026-06-21. [Online]. Available: <https://epoch.ai/data-insights/hyperscaler-capex-vs-cash-flow>.
- [6] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024. URL <https://arxiv.org/abs/2407.21783>.
- [7] Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, et al. Mixtral of experts. *arXiv preprint arXiv:2401.04088*, 2024. URL <https://arxiv.org/abs/2401.04088>.
- [8] Alexey Dosovitskiy, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, et al. An image is worth 16x16 words: Transformers for image recognition at scale. In *Proc. ICLR*, 2021. URL <https://arxiv.org/abs/2010.11929>.
- [9] Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. BLIP-2: Bootstrapping language-image pre-training with frozen image encoders and large language models. In *Proc. ICML*, 2023. URL <https://arxiv.org/abs/2301.12597>.
- [10] Jeff Johnson, Matthijs Douze, and Hervé Jégou. Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3):535–547, 2019. URL <https://arxiv.org/abs/1702.08734>.
- [11] Zixuan Zhou, Xuefei Ning, Ke Hong, Tianyu Fu, Jiaming Xu, Shiyao Li, Yuming Lou, Luning Wang, Zhihang Yuan, Xiuhong Li, Shengen Yan, Guohao Dai, Xiaoping Zhang, Yuhan Dong, and Yu Wang. A survey on efficient inference for large language models. *arXiv preprint arXiv:2404.14294*, 2024. URL <https://arxiv.org/abs/2404.14294>.
- [12] Li Baolin et al. LLM inference serving: Survey of recent advances and opportunities. In *Proc. HPEC*, 2024. URL <https://arxiv.org/abs/2407.12391>.
- [13] Miao Xupeng et al. Towards efficient generative large language model serving: A survey from algorithms to systems. *ACM Computing Surveys*, 58(1):1–37, 2025. URL <https://arxiv.org/abs/2312.15234>.
- [14] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. *arXiv preprint arXiv:2309.06180*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- [15] Gyeong-In Yu, Joo Seong Jeong, Geon-Woo Kim, Soojeong Kim, and Byung-Gon Chun. Orca: A distributed serving system for transformer-based generative models. In *Proc. OSDI*, pages 521–538, 2022. URL <https://www.usenix.org/>

- conference/osdi22/presentation/you.
- [16] Du Kuntai et al. Prefillonly: An inference engine for prefill-only workloads in large language model applications. In *Proc. SOSP*, 2025. URL <https://arxiv.org/abs/2505.07203>.
- [17] Jin Yibo et al. P/d-serve: Serving disaggregated large language model at scale. *arXiv preprint arXiv:2408.08147*, 2024. URL <https://arxiv.org/abs/2408.08147>.
- [18] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Daniel Y. Fu, Zhiqiang Xie, Beidi Chen, Clark Barrett, Joseph E. Gonzalez, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. Flexgen: High-throughput generative inference of large language models with a single gpu. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2303.06865>.
- [19] Reza Yazdani Aminabadi, Samyam Rajbhandari, Ammar Ahmad Awan, Chenyang Li, Dong Li, Erich Zheng, Olatunji Ruwase, and Yuxiong He. Deep-speed inference: Enabling efficient inference of transformer models at unprecedented scale. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*, 2022. URL <https://arxiv.org/abs/2207.00032>.
- [20] Xu Yi et al. Pie: Pooling CPU memory for LLM inference. *arXiv preprint arXiv:2411.09317*, 2024. URL <https://arxiv.org/abs/2411.09317>.
- [21] Xuanlin Jiang et al. Neo: Saving GPU memory crisis with CPU offloading for online llm inference. *arXiv preprint arXiv:2411.01142*, 2024. URL <https://arxiv.org/abs/2411.01142>.
- [22] Xiangwen Zhuge et al. SpecOffload: Unlocking latent GPU capacity for LLM inference on resource-constrained devices. *arXiv preprint arXiv:2505.10259*, 2025. URL <https://arxiv.org/abs/2505.10259>.
- [23] Xu Jiale et al. eLLM: Elastic memory management framework for efficient LLM serving. *arXiv preprint arXiv:2506.15155*, 2025. URL <https://arxiv.org/abs/2506.15155>.
- [24] Chien-Chin Huang, Gu Jin, and Jinyang Li. Swapadvisor: Pushing deep learning beyond the gpu memory limit via smart swapping. In *Proceedings of the 25th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 2020. URL <https://doi.org/10.1145/3373376.3378530>.
- [25] Wang Kun et al. Swapnet: Efficient swapping for dnn inference on edge ai devices beyond the memory budget. *IEEE Transactions on Mobile Computing*, 23(9): 8935–8950, 2024. URL <https://arxiv.org/abs/2401.16757>.
- [26] Kida Shun et al. Revisiting memory swapping for big-memory applications. In *Proc. HPCASIA*, 2025. URL <https://doi.org/10.1145/3712031.3712332>.
- [27] Elias Frantar and Dan Alistarh. SparseGPT: Massive language models can be accurately pruned in one-shot. In *Proc. ICML*, 2023. URL <https://arxiv.org/abs/2301.00774>.
- [28] Tri Dao, Daniel Y. Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and Memory-Efficient Exact Attention with IO-Awareness. In *Proc. NeurIPS*, 2022. URL <https://arxiv.org/abs/2205.14135>.

- [29] Korthikanti Vijay Anand et al. Reducing activation recomputation in large transformer models. *Proc. MLSys*, 5:341–353, 2023. URL <https://arxiv.org/abs/2205.05198>.
- [30] Huang Yanpin et al. GPipe: Efficient training of giant neural networks using pipeline parallelism. In *Proc. NeurIPS*, 2019. URL <https://arxiv.org/abs/1811.06965>.
- [31] Mohammad Shoeybi, Mostofa Patwary, Raul Puri, Patrick LeGresley, Jared Casper, and Bryan Catanzaro. Megatron-LM: Training multi-billion parameter language models using model parallelism. *arXiv preprint arXiv:1909.08053*, 2019. URL <https://arxiv.org/abs/1909.08053>.
- [32] A.H. Hunter et al. Beyond malloc efficiency to fleet efficiency: a hugepage-aware memory allocator. In *Proc. OSDI*, 2021. URL <https://www.usenix.org/conference/osdi21/presentation/hunter>.
- [33] Yang Hanmei et al. Numalloc: A faster numa memory allocator. In *Proc. ISMM*, 2023. URL <https://doi.org/10.1145/3591195.3595276>.
- [34] Ugljesa Milic, Oreste Villa, Evgeny Bolotin, Akhil Arunkumar, Eiman Ebrahimi, Aamer Jaleel, Alex Ramirez, and David Nellans. Beyond the socket: NUMA-aware GPUs. In *Proc. MICRO*, pages 123–135, 2017. URL <https://doi.org/10.1145/3123939.3124534>.
- [35] Carl A. Waldspurger. Memory resource management in VMware ESX Server. In *Proceedings of the 5th USENIX Symposium on Operating Systems Design and Implementation (OSDI)*, pages 181–194, 2002. URL <https://www.usenix.org/conference/osdi-02/memory-resource-management-vmware-esx-server>.
- [36] Long Sifan et al. A survey on intelligent network operations and performance optimization based on large language models. *IEEE Communications Surveys & Tutorials*, 2025. URL <https://ieeexplore.ieee.org/document/10829820>.
- [37] Zhi Wang et al. Model offloading techniques for edge AI: A survey. *IEEE Internet of Things Journal*, 2023. URL <https://doi.org/10.1109/JIOT.2023.3254417>. DOI: 10.1109/JIOT.2023.3254417.
- [38] Sarkar Rishov et al. Edge-moe: Memory-efficient multi-task vision transformer architecture with task-level sparsity via mixture-of-experts. In *Proc. ICCAD*, 2023. URL <https://arxiv.org/abs/2305.18691>.
- [39] Luo Haoxiang et al. Toward edge general intelligence with multiple-large language model (Multi-LLM): architecture, trust, and orchestration. *IEEE Trans. Cogn. Commun. Netw.*, 2025. URL <https://arxiv.org/abs/2507.00672>.
- [40] He Zixuan et al. CPU–GPU heterogeneous computation offloading and resource allocation scheme for industrial internet of things. *IEEE Internet of Things Journal*, 11(6):11152–11164, 2023. URL <https://ieeexplore.ieee.org/document/10319405>.
- [41] Beaumont Olivier et al. Optimal GPU-CPU offloading strategies for deep neural network training. In *Proc. Euro-Par*, 2020. URL https://doi.org/10.1007/978-3-030-57675-2_10.
- [42] R. L. Mattson, J. Gecsei, D. R. Slutz, and I. L. Traiger. Evaluation techniques for storage hierarchies. *IBM Systems Journal*, 9(2):78–117, 1970. doi: 10.1147/sj.92.0078. URL <https://doi.org/10.1147/sj.92.0078>.

- 428 [43] Samuel Williams, Andrew Waterman, and David Patterson. Roofline: an insight-
429 ful visual performance model for multicore architectures. *Communications of the*
430 *ACM*, 52(4):65–76, 2009. doi: 10.1145/1498765.1498785. URL [https://doi.org/](https://doi.org/10.1145/1498765.1498785)
431 [10.1145/1498765.1498785](https://doi.org/10.1145/1498765.1498785).
- 432 [44] Leonard Kleinrock. *Queueing Systems, Volume 1: Theory*. John Wiley & Sons,
433 New York, 1975. URL <https://dl.acm.org/doi/book/10.5555/1096491>. ISBN 978-
434 0-471-49110-1.