

Two dimensionless ratios define operational limits in hierarchical-memory inference

Geunsik Lim ^{1*}

¹ College of Information and Communication Engineering,
Sungkyunkwan University (SKKU), Suwon, South Korea .

Corresponding author(s). E-mail(s): leemgs@g.skku.edu;

Abstract

We formulate hierarchical-memory inference using two dimensionless ratios: fast-memory residency $R_C = C_{\text{fast}}/W$ and compute-transfer overlap $R_B = T_{\text{comp}}/T_{\text{transfer}}$. These ratios define a capacity-limited region ($R_C < 0.5$), an I/O-limited region ($R_C \geq 0.5$, $R_B < 1$), and a coordination-dominated region ($R_C \geq 0.5$, $R_B \geq 1$). We derive irreducible capacity and transfer-cost bounds and release a measurement harness using wall-clock or device-event timing. A compiled dependent-load microbenchmark on one Intel Xeon CPU shows mean access time increasing from 6.8 ns for a 0.5 MiB working set to 169.7 ns for 256 MiB. The increase spans several cache and translation levels and does not identify a discontinuity at $R_C = 0.5$. On an NVIDIA Tesla T4 GPU, the same harness confirms both falsifiable predictions: per-step latency rises 2.83-fold as residency falls below $R_C = 0.5$, and the I/O-limited label appears exactly where the derived overlap boundary $R_B < 1$ predicts. These results establish the two-ratio labels as a measurement discipline on real silicon; broad accelerator and production-workload generality remains open.

Keywords: hierarchical memory systems, memory orchestration, AI inference, working-set residency, compute-transfer overlap, reproducible measurement

1 Introduction

Modern foundation models have shifted the performance frontier of AI inference from arithmetic throughput to *data coordination*. When model weights, activations, and key-value (KV) caches exceed accelerator memory, inference latency is governed less by compute and more by how efficiently execution coordinates data movement across

a hierarchical memory system spanning device memory, host memory, and secondary storage [1, 2].

Hardware analyses identify memory capacity, bandwidth, and interconnect latency as important constraints on inference [3]. Separating capacity pressure from exposed transfer is therefore useful when reporting and comparing inference systems.

Existing systems reduce memory pressure via CPU/NVMe offloading [4, 5], GPU–CPU swapping [6], and KV-cache management [7]. Their outcomes depend on model size, batch size, hardware topology, and the amount of traffic that remains on the critical path. This motivates a compact description that exposes those conditions before comparing policies.

We take a regime-based view. Denning’s working-set model and thrashing theory [8] identify capacity saturation as a binary event on a *single-tier* system, and the Roofline model [9] expresses performance as $\min(F \cdot I, B_{\text{peak}})$ for a fixed workload on a single tier. Neither framework addresses a *multi-tier* hierarchy in which computation, locality, cross-tier transfer, and synchronisation co-exist as distinct, adjustable latency terms. Our accounting model decomposes end-to-end latency into four terms ($T_{\text{comp}} + T_{\text{mem}} + T_{\text{swap}} + T_{\text{sync}}$) and uses two dimensionless ratios to label an operating point. The three operational labels are governed by the fast-memory residency ratio $R_C = C_{\text{fast}}/W$ and the overlap ratio $R_B = B_{\text{slow}}T_{\text{comp}}/D = T_{\text{comp}}/T_{\text{transfer}}$. The labels state whether majority residency is available and whether ideal compute–transfer overlap can hide compulsory traffic; they do not by themselves predict a policy’s speed-up.

We make three contributions:

- **Operational formulation.** We separate residency from compute–transfer overlap using two measurable, dimensionless ratios and state an explicit three-label classification rule.
- **Minimal analytical model.** A structural lower bound on the irreducible memory-access and transfer costs, with regime boundaries following directly from the two ratios: $\theta_B = 1.0$ derived as the overlap boundary ($R_B = T_{\text{comp}}/T_{\text{transfer}} = 1$), and $\theta_C = 0.50$ as the majority-residency crossover.
- **Open proof of concept.** A compiled dependent-load CPU probe illustrates working-set sensitivity, and a layered-matrix harness runs on both a CPU and a single NVIDIA GPU (Tesla T4), where the residency and overlap predictions both hold. We distinguish this single-accelerator proof of concept, and the unexercised XLA/TPU path, from a multi-platform campaign.

Together, the formulation and measurements provide a reproducible starting point for testing when hierarchical-memory coordination is constrained by residency or exposed transfer. They confirm the two labels on one CPU and one GPU but do not establish effect sizes, strategy rankings, or cross-accelerator generality.

2 Results

2.1 A two-ratio operational model

We describe a hierarchical-memory operating point using

$$R_C = \frac{C_{\text{fast}}}{W}, \quad R_B = \frac{B_{\text{slow}} T_{\text{comp}}}{D} = \frac{T_{\text{comp}}}{T_{\text{transfer}}}, \quad (1)$$

where C_{fast} is usable fast-memory capacity, W is the active working set, D is compulsory cross-tier traffic per step, B_{slow} is sustained bandwidth on the limiting link, and $T_{\text{transfer}} = D/B_{\text{slow}}$. Unlike a ratio that uses total step time, R_B is not circular: it compares independently timed computation and transfer intervals.

We use $\theta_C = 0.50$ as an explicit majority-residency convention and derive $\theta_B = 1$ from equality of computation and transfer time. These definitions produce three operational labels:

- **capacity-limited:** $R_C < \theta_C$;
- **I/O-limited:** $R_C \geq \theta_C$ and $R_B < \theta_B$; and
- **coordination-dominated:** $R_C \geq \theta_C$ and $R_B \geq \theta_B$.

The capacity boundary is a convention, not a thermodynamic critical point. The I/O boundary is an overlap boundary: for $R_B < 1$, at least $T_{\text{transfer}} - T_{\text{comp}}$ cannot be hidden by ideal overlap.

2.2 Irreducible costs and falsifiable predictions

For accounting purposes, define non-overlapping critical-path contributions so that a step can be written as

$$T_{\text{total}} = T_{\text{comp}} + T_{\text{mem}} + T_{\text{swap}} + T_{\text{sync}}. \quad (2)$$

Independently of that accounting decomposition, a conservative service-time lower bound that allows the underlying services to overlap is

$$T_{\text{total}} \geq \max \left\{ T_{\text{comp}}, \rho W \max(0, 1 - R_C), \frac{D}{B_{\text{slow}}} \right\}, \quad (3)$$

where W denotes bytes accessed in the step and ρ is a lower bound on service time per non-resident byte for the specified access pattern. Taking the maximum permits ideal overlap and avoids double-counting bytes that contribute to both a cache miss and a cross-tier transfer. Synchronisation and contention can only increase observed time, so Equation (3) is not a performance predictor. The residual above it cannot be assigned to one latency component without event-level measurements.

The formulation makes two falsifiable predictions. First, reducing residency while holding the access pattern fixed should expose a higher memory-access cost. Second, when $R_C \geq \theta_C$, configurations with $R_B < 1$ must expose transfer time even under ideal compute-transfer overlap. It does not, without additional measurements, predict




Fig. 1 Analytical regime map. The vertical line is the declared majority-residency convention $\theta_C = 0.50$; the horizontal segment is the derived overlap boundary $\theta_B = 1.0$. The diagram is generated from the same constants as the released classifier and contains no empirical data.

the magnitude of a strategy’s speed-up or claim that strategy rankings necessarily invert.

2.3 CPU proof-of-concept measurements

We evaluated working-set sensitivity on an Intel Xeon Platinum 8370C CPU exposed through a virtualised Linux environment. The experiment used a randomised pointer chain over working sets from 0.5 to 256 MiB and seven trials per point. A compiled C loop performed dependent loads and timed them with `CLOCK_MONOTONIC_RAW`; Python was used only to construct the chain and summarise trials. Linux `sysfs` reported a 48 MiB last-level cache, which we used as C_{fast} . Mean dependent-load time was 6.8 ± 0.1 ns at 0.5 MiB ($R_C = 96$), 118.5 ± 6.6 ns at 48 MiB ($R_C = 1$), 132.5 ± 14.3 ns at 96 MiB ($R_C = 0.5$), and 169.7 ± 10.4 ns at 256 MiB ($R_C = 0.188$), where uncertainty is one standard deviation across trials. The non-monotonic intermediate points show that cache level, translation, virtualisation, and system noise matter. In particular, these data do not show a discontinuity at the conventional $R_C = 0.5$ label.

As a separate harness check, we used the layered matrix program with 12 layers, $d = 512$, batch size 8, and five timing windows. It varied the fraction of layer weights retained in the fast tier. The mean total latency averaged over the three sampled points below $R_C = 0.5$ was 4.03 ms, compared with 1.87 ms for the fully resident point, a 2.15-fold descriptive ratio. Because computation and copying share the same CPU

Table 1 Selected measured points from the compiled CPU pointer-chain experiment. Values are mean $\pm$ one standard deviation over seven trials.

Working set (MiB)	R_C	Latency (ns/load)	Interpretation
0.5	96.000	6.8 ± 0.1	small working set
48	1.000	118.5 ± 6.6	reported LLC capacity
96	0.500	132.5 ± 14.3	residency convention
256	0.188	169.7 ± 10.4	largest working set

memory system, the accompanying R_B sweep was noisy and is not used as evidence for the I/O boundary.

These measurements are proof-of-concept results from one CPU, not a substitute for the previously planned A100, MI250, TPU v4, Inferentia2, and Optane campaign. Accordingly, we make no quantitative cross-accelerator, classifier- accuracy, energy, or strategy-inversion claim in this version.

2.4 GPU proof-of-concept measurements

We repeated the layered program on a real accelerator—an NVIDIA Tesla T4 (Google Colab)—with $d = 2048$, 24 layers, batch size 8, and ten timing windows, using CUDA events to time computation and host-to-device transfer separately. Non-resident layer weights were staged in pinned host memory and copied per step, so residency (R_C) and overlap (R_B) were varied by construction. This run tests the two falsifiable predictions of Section 2.2 on hardware whose device–host link, unlike a single CPU’s memory system, makes the exposed-transfer regime physically reachable.

Prediction 1 (residency). Reducing residency raised per-step latency monotonically, from 14.53 ms at full residency ($R_C = 1$) to a mean of 41.08 ms across the capacity-limited points ($R_C < \theta_C$), a 2.83-fold increase—the higher memory-access cost that Prediction 1 requires.

Prediction 2 (overlap). Holding residency fixed and scaling compute per layer moved the operating point across the derived boundary $\theta_B = 1$: every point with $R_B < 1$ was labelled I/O-limited and every point with $R_B \geq 1$ coordination-dominated, with the transition falling between $R_B = 0.90$ and $R_B = 1.19$ as predicted, not at a fitted value. Total latency over the exposed-transfer points averaged 23.87 ms; because these points also carry the smallest compute, the comparison with the overlapped points (35.95 ms) is confounded by compute scaling and is not read as an inversion. Unlike the CPU harness, this GPU run therefore does provide direct evidence for the overlap boundary.

This remains a single-accelerator, proof-of-concept result; it confirms the two qualitative predictions on real GPU silicon but does not establish cross-accelerator boundary values, strategy inversion, classifier accuracy, or energy claims, which still require the full multi-platform campaign.

2.5 Reproducibility boundary

The repository contains the raw JSONL records, machine-readable summaries, and scripts used for Table 1. It also contains a simulated backend for testing the analysis pipeline. Simulated output is labelled as such and is not used as empirical evidence. Accelerator support in the measurement harness is an executable protocol for future validation; it is not evidence that those accelerators were measured in this study.

3 Discussion

The two-ratio description separates two questions that are often conflated in hierarchical-memory optimisation: whether the active data can reside in the fast tier, and whether compulsory transfer can be hidden behind computation. The resulting labels are operational rather than universal phases. $\theta_B = 1$ follows exactly from the overlap definition, whereas $\theta_C = 0.5$ is a declared midpoint that makes “majority resident” precise. A different application may choose another residency threshold without changing Equation (3).

The CPU probes show that latency depends strongly on working-set residency, but they do not validate a sharp boundary at $R_C = 0.5$. They stress different mechanisms: dependent pointer loads probe cache and translation effects, while the layered matrix harness combines weight copying with computation. The single-GPU run (Section 2.4) does exercise accelerator DMA and asynchronous CUDA streams, and it confirms both qualitative predictions, including the overlap crossover at $\theta_B = 1$ that the CPU harness could not resolve. It nonetheless remains one accelerator: the data do not establish cross-accelerator boundary values, HBM- or collective-communication effects at scale, or the generality of the three labels across model families.

The framework is most useful as a measurement discipline. Reporting R_C and R_B alongside latency makes hidden assumptions about residency and overlap inspectable. It also prevents a common circularity: using total step duration in both the numerator of a bandwidth ratio and as the outcome being explained. For deployment, C_{fast} , W , D , sustained bandwidth, and isolated compute time must be measured for the actual workload and hardware rather than copied from nominal data-sheet values.

Several limitations are material. The present evidence comes from one CPU and one GPU with small synthetic probes; it does not include production models, concurrent requests, accelerator energy measurements, strategy comparisons, or unseen-hardware classifier tests. The standard deviations describe repeated trials on one machine and are not confidence intervals over a population of platforms. These measurements are exploratory technical checks rather than independent hardware replicates, and no causal inference is supported. The pointer-chain endpoints also span several hierarchy transitions, so their ratio should not be interpreted as a universal effect size. Finally, Equation (2) is an accounting identity only when its terms are defined as non-overlapping critical-path contributions; the max-form lower bound avoids assuming that decomposition in measured systems.

A decisive follow-up should preregister operating-point grids around $R_C = 0.5$ and $R_B = 1$, collect raw event traces on at least two accelerator architectures, and compare fixed orchestration policies without changing other workload parameters. Such

193 data could test boundary sharpness, policy ranking, and generalisation. Until those
 194 measurements are available, the contribution of this work is the corrected analytical
 195 formulation, the open protocol, and a CPU-and-single-GPU proof of concept rather
 196 than a claim of universal accelerator validation.

197 4 Methods

198 Definitions and classification

199 We calculate $R_C = C_{\text{fast}}/W$ from a declared usable fast-tier capacity and active
 200 working-set size. We calculate $R_B = T_{\text{comp}}/T_{\text{transfer}}$, with compute and transfer timed
 201 separately. Classification gives capacity precedence: $R_C < 0.5$ is capacity-limited;
 202 otherwise $R_B < 1$ is I/O-limited; all remaining points are coordination-dominated.
 203 This rule is implemented in the released Python code.

204 Pointer-chain capacity probe

205 The CPU probe allocates a NumPy array and constructs a seeded random permu-
 206 tation. A C11 kernel compiled with `-O3` traverses it as a dependent pointer chain;
 207 each loaded index determines the address of the next load, limiting vectorisation,
 208 instruction-level parallelism, and prefetching. We measured working sets of 0.5, 1, 2, 4,
 209 8, 16, 24, 32, 48, 64, 96, 128, 192, and 256 MiB. Each point used seven timed trials after
 210 warm-up. The output stores working-set size, R_C , mean nanoseconds per dependent
 211 load, standard deviation, and trial count. The fast-tier capacity was read from Linux
 212 `sysfs` and was 48 MiB. The host exposed three logical CPUs of an Intel Xeon Platinum
 213 8370C under KVM; the summary records the kernel, Python, and NumPy versions.
 214 Timing used `clock_gettime(CLOCK_MONOTONIC_RAW)` inside the compiled loop. No
 215 outlier was removed and CPU affinity or frequency was not controlled.

216 Layered matrix and copy probe

217 The second probe allocates square, float32 layer matrices; a selected fraction remains
 218 resident and non-resident matrices are copied before multiplication. Matrix entries
 219 are deterministically seeded and scaled by $1/\sqrt{d}$ to prevent numerical overflow during
 220 repeated multiplication. On the CPU we used NumPy’s wall-clock path with 12 layers,
 221 $d = 512$, and five windows per point. On the GPU (Section 2.4) the same program
 222 ran through the harness’s CUDA-event path on an NVIDIA Tesla T4 with 24 layers,
 223 $d = 2048$, batch size 8, ten windows, and pinned host staging; computation and
 224 host-to-device transfer were timed separately with `torch.cuda.Event`. The harness
 225 additionally supports XLA synchronisation for TPUs, which was not exercised here.

226 For the reported capacity ratio, we averaged total latency at the three sampled
 227 points below $R_C = 0.5$ and divided by latency at $R_C = 1$. This endpoint summary is
 228 descriptive; no hypothesis test or population-level confidence interval is claimed. The
 229 pointer-chain table reports mean and population standard deviation across the seven
 230 repeated trials exactly as stored in the JSON summary.

231 **Software and provenance**

232 The compiled pointer-chain run used Python 3.12.13, NumPy 2.5.2, and GCC 13.3;
233 the earlier layered-matrix harness check used Python 3.11.15, but that summary did
234 not capture the NumPy version. Seeds, grids, timing code, raw records, summaries,
235 and commands are version controlled. The analytical regime map reads threshold
236 constants from the same module as the classifier. The simulated backend is provided
237 solely for software testing and qualitative exploration; its outputs are excluded from
238 the reported measurements.

239 **Statistics and reporting**

240 This proof-of-concept study did not use random assignment, blinding, or a sample-size
241 calculation. Trial counts were fixed before summarisation. We report all measured grid
242 points in the machine-readable data, selected points in Table 1, arithmetic means, and
243 one standard deviation. No claim of statistical significance is made. Hardware-platform
244 generalisation requires independent machines and is left for future work.

245 **Data availability**

246 All data supporting the reported CPU measurements are included in the pub-
247 lic repository at <https://github.com/leemgs/orion>, under `code/results/cpu_probe/`
248 and `code/results/colab_probe/`. These directories contain the raw JSONL records
249 and machine-readable JSON summaries used in the text and table. No accelerator
250 data are reported in this article.

251 **Code availability**

252 The analysis and measurement code is available at <https://github.com/leemgs/orion>.
253 The `code/experiments/` directory contains the CPU probes, analytical-figure gener-
254 ator, and optional CUDA/XLA measurement paths. Exact reproduction commands
255 are provided in the README and Supplementary Information. A separately labelled
256 simulated backend is supplied for software testing; simulated output is not used as
257 evidence in this article.

258 **Author Contributions**

259 G.L.: Conceptualization, Formal analysis, Investigation, Methodology, Software, Val-
260 idation, Visualization, Writing – original draft, Writing – review & editing.

261 **Competing Interests**

262 The author declares no competing interests.

Acknowledgements

We thank Donghyeon Kang, Seungmin Oh and Yunsu Lee for valuable feedback. We are also grateful to Wooyeon Lee, Hyoyoung Kim, and Heekuk Lee, experts in systems and memory, whose insights and meaningful comments greatly improved this work.

References

- [1] Xiaoyu Ma and David Patterson. Challenges and research directions for large language model inference hardware. *IEEE Computer*, 59, 2026. URL <https://arxiv.org/abs/2601.05047>. Accepted.
- [2] Amir Gholami, Zhewei Yao, Sehoon Kim, Coleman Hooper, Michael W. Mahoney, and Kurt Keutzer. Ai and memory wall. *IEEE Micro*, 2024. URL <https://arxiv.org/abs/2403.14123>.
- [3] Zixuan Zhou, Xuefei Ning, Ke Hong, Tianyu Fu, Jiaming Xu, Shiyao Li, Yuming Lou, Luning Wang, Zhihang Yuan, Xiuhong Li, Shengen Yan, Guohao Dai, Xiaoping Zhang, Yuhang Dong, and Yu Wang. A survey on efficient inference for large language models. *arXiv preprint arXiv:2404.14294*, 2024. URL <https://arxiv.org/abs/2404.14294>.
- [4] Ying Sheng, Lianmin Zheng, Binhang Yuan, Zhuohan Li, Max Ryabinin, Daniel Y. Fu, Zhiqiang Xie, Beidi Chen, Clark Barrett, Joseph E. Gonzalez, Percy Liang, Christopher Ré, Ion Stoica, and Ce Zhang. Flexgen: High-throughput generative inference of large language models with a single gpu. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023. URL <https://arxiv.org/abs/2303.06865>.
- [5] Reza Yazdani Aminabadi, Samyam Rajbhandari, Ammar Ahmad Awan, Chenyang Li, Dong Li, Erich Zheng, Olatunji Ruwase, and Yuxiong He. Deepspeed inference: Enabling efficient inference of transformer models at unprecedented scale. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis (SC)*, 2022. URL <https://arxiv.org/abs/2207.00032>.
- [6] Chien-Chin Huang, Gu Jin, and Jinyang Li. Swapadvisor: Pushing deep learning beyond the gpu memory limit via smart swapping. In *Proceedings of the 25th International Conference on Architectural Support for Programming Languages and Operating Systems (ASPLOS)*, 2020. URL <https://doi.org/10.1145/3373376.3378530>.
- [7] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. *arXiv preprint arXiv:2309.06180*, 2023. URL <https://arxiv.org/abs/2309.06180>.
- [8] R. L. Mattson, J. Gecsei, D. R. Slutz, and I. L. Traiger. Evaluation techniques for storage hierarchies. *IBM Systems Journal*, 9(2):78–117, 1970. doi: 10.1147/sj.92.0078. URL <https://doi.org/10.1147/sj.92.0078>.
- [9] Samuel Williams, Andrew Waterman, and David Patterson. Roofline: an insightful visual performance model for multicore architectures. *Communications of the*

305 *ACM*, 52(4):65–76, 2009. doi: 10.1145/1498765.1498785. URL [https://doi.org/10.](https://doi.org/10.1145/1498765.1498785)
306 [1145/1498765.1498785](https://doi.org/10.1145/1498765.1498785).